

# Paper Evaluation

July 14, 2020

```
[1]: import os
import shutil

import pandas as pd
import logpred_method as experiment

from sklearn.model_selection import train_test_split

# Use "FRACTION = None" for full dataset
FRACTION: float = None

# lr: Linear Regression
# ab: Ada Boost
# rf: Random Forest
# dt: Decision Tree
# et: Extra Trees
MODELS = ["lr", "ab", "rf", "dt", "et"]

# You can ignore features on the experiment
IGNORED_FEATURES = ["tryCatchQty_class", "tryCatchQty_method"]

# Hyperparameter tuning
TUNING_ENABLED = True

# Stores estimators and feature importances across experiments
ESTIMATORS = {}
FEATURE_IMPORTANCES = {}
```

## 1 Utilities

```
[2]: def merge_scores(scores):  
    """  
    Returns a merged score from a sequence of scores.  
    This is useful to see scores as Pandas DataFrames.  
  
    Example:  
    in - [{"a": 1, "b": 2}, {"a": 10, "b": 20}]  
    out - {"a": [1, 10], "b": [2, 20]}  
    """  
    merged = {k: [] for k in scores[0].keys()}  
    for score in scores:  
        for k, v in score.items():  
            merged[k].append(v)  
  
    return merged
```

## 2 Experiment CSV and Output directory

```
[3]: csv_path = os.path.abspath(os.path.join("out", "dataset", "adyen-main",  
    ↪ "dataset_full.csv"))  
  
X_adyen, y_adyen = experiment.load_dataset(csv_path, drops=IGNORED_FEATURES,  
    ↪ fraction=FRACTION)  
X_adyen_train, X_adyen_test, y_adyen_train, y_adyen_test = train_test_split(  
    X_adyen, y_adyen, test_size=0.2, stratify=y_adyen, random_state=experiment.  
    ↪ RANDOM_SEED  
)  
  
output_dir = os.path.abspath(os.path.join("out", "ml",  
    ↪ f"evaluation-tuning-{TUNING_ENABLED}"))  
if os.path.exists(output_dir):  
    shutil.rmtree(output_dir)  
os.makedirs(output_dir)
```

## 3 RQ 1. What is the performance of machine learning models in predicting log placement in a large-scale enterprise system?

```
[4]: def rq1():  
    scores = []  
    for model in MODELS:  
        out = experiment.run(  
            model,  
            X_train=X_adyen_train,
```

```

        X_test=X_adyen_test,
        y_train=y_adyen_train,
        y_test=y_adyen_test,
        output_to=os.path.join(output_dir, f"rq1-{model}.log"),
        tuning_enabled=TUNING_ENABLED
    )
    estimator, score, fi = out
    scores.append(score)

    # Save to the global state this run
    ESTIMATORS[model] = estimator
    FEATURE_IMPORTANCES[model] = fi

    return scores

rq1_scores = rq1()

```

### 3.1 Results

```

[5]: results_rq1 = pd.DataFrame.from_dict(merge_scores(rq1_scores)).
    ↪set_index(["model"])
    results_rq1.reset_index().to_csv(
        os.path.join(output_dir, "rq1-results.csv"),
        index=False,
    )
    results_rq1["prec recall acc tn fp fn tp total".split(" ")]

```

```

[5]:
      prec  recall  acc  tn  fp  fn  tp  total
model
lr      0.656053  0.373446  0.678597  56232  929  2973  1772  61906
ab      0.645349  0.444468  0.712096  56002  1159  2636  2109  61906
rf      0.814496  0.618124  0.803219  56493   668  1812  2933  61906
dt      0.585751  0.523288  0.746284  55405  1756  2262  2483  61906
et      0.740093  0.570706  0.777034  56210   951  2037  2708  61906

```

## 4 RQ 2. What is the impact of different class balancing strategies on prediction?

```

[6]: # Similar to rq1 but we include sampling in the experiment now.
def rq2():
    scores = []
    for model in MODELS:
        for balancing in ["smote", "rus"]:
            out = experiment.run(
                model,

```

```

        X_train=X_adyen_train,
        X_test=X_adyen_test,
        y_train=y_adyen_train,
        y_test=y_adyen_test,
        balancing=balancing,
        output_to=os.path.join(output_dir, f"rq2-{model}-{balancing}.
→log"),
        tuning_enabled=TUNING_ENABLED
    )
    estimator, score, fi = out
    scores.append(score)

    # Save to the global state this run
    key = f"{model}-{balancing}"
    ESTIMATORS[key] = estimator
    FEATURE_IMPORTANCES[key] = fi

    return scores

rq2_scores = rq2()

```

## 4.1 Results

```

[7]: results_rq2 = pd.DataFrame.from_dict(merge_scores(rq2_scores)).
→set_index(["model", "balancing"])
results_rq2.reset_index().to_csv(
    os.path.join(output_dir, "rq2-results.csv"),
    index=False,
)

relevant_cols = "prec recall acc".split(" ")
results_rq2[relevant_cols]

```

```

[7]:

```

		prec	recall	acc
model	balancing			
lr	smote	0.385021	0.891675	0.886724
	rus	0.382456	0.895890	0.887904
ab	smote	0.314637	0.948156	0.888355
	rus	0.371991	0.931507	0.900482
rf	smote	0.504170	0.891886	0.909537
	rus	0.412009	0.961644	0.923860
dt	smote	0.389187	0.872287	0.879321
	rus	0.323624	0.953003	0.893832
et	smote	0.457025	0.907692	0.909087
	rus	0.408215	0.950896	0.918232

Comparative result to the baseline (no balancing). Positive value indicates improvement.

```
[8]: results_rq2_rel = results_rq2.loc[MODELS, relevant_cols] - results_rq1.  
      ↪loc[MODELS, relevant_cols]  
      results_rq2_rel.reset_index().to_csv(  
          os.path.join(output_dir, "rq2-results-relative.csv"),  
          index=False  
      )  
      results_rq2_rel
```

```
[8]:
```

		prec	recall	acc
model	balancing			
lr	smote	-0.271032	0.518230	0.208127
	rus	-0.273597	0.522445	0.209308
ab	smote	-0.330711	0.503688	0.176259
	rus	-0.273358	0.487039	0.188386
rf	smote	-0.310326	0.273762	0.106318
	rus	-0.402487	0.343519	0.120641
dt	smote	-0.196565	0.348999	0.133038
	rus	-0.262127	0.429715	0.147548
et	smote	-0.283068	0.336986	0.132052
	rus	-0.331878	0.380190	0.141198

## 5 RQ 3. How do machine learning models perceive predictors?

```
[9]: def rank_to_df(rank):  
      return pd.DataFrame.from_records(  
          [(name, sum(count), *count) for name, count in rank.items()],  
          columns="feature total 1st 2nd 3rd".split(" "),  
          ).sort_values(by="total 1st 2nd 3rd".split(" "), ascending=False)  
  
      def feature_importance_rank(selected_models, top_n=3):  
          rank = {}  
          for model in selected_models:  
              ordered_features = sorted(  
                  FEATURE_IMPORTANCES[model],  
                  key=lambda pair: pair[1],  
                  reverse=True  
              )  
              for pos, feature_pair, in enumerate(ordered_features[:top_n]):  
                  feature = feature_pair[0]  
                  if feature not in rank.keys():  
                      rank[feature] = [0 for i in range(top_n)]  
                  rank[feature][pos] += 1  
          return rank
```

## 5.1 Results

```
[10]: fi = rank_to_df(feature_importance_rank(MODELS))
      fi.to_csv(
          os.path.join(output_dir, "rq3-fi-regular.csv"),
          index=False
      )
      fi
```

```
[10]:
```

	feature	total	1st	2nd	3rd
0	maxNestedBlocks	4	4	0	0
3	loc_method	2	1	0	1
7	cbo_method	2	0	2	0
2	uniqueWordsQty_method	2	0	0	2
1	maxNestedBlocksQty	1	0	1	0
4	cbo_class	1	0	1	0
6	wmc_method	1	0	1	0
5	publicMethodsQty	1	0	0	1
8	loopQty_method	1	0	0	1

```
[11]: fi_smote = rank_to_df(
      feature_importance_rank([
          model_key
          for model_key in FEATURE_IMPORTANCES.keys()
          if "smote" in model_key
      ])
      )
      fi_smote.to_csv(
          os.path.join(output_dir, "rq3-fi-smote.csv"),
          index=False
      )
      fi_smote
```

```
[11]:
```

	feature	total	1st	2nd	3rd
1	maxNestedBlocks	5	4	1	0
5	wmc_method	2	0	2	0
2	cbo_method	2	0	0	2
0	loc_method	1	1	0	0
3	constructor_False	1	0	1	0
8	uniqueWordsQty_method	1	0	1	0
4	constructor_True	1	0	0	1
6	rfc_method	1	0	0	1
7	variablesQty_method	1	0	0	1

```
[12]: fi_rus = rank_to_df(
      feature_importance_rank([
          model_key
          for model_key in FEATURE_IMPORTANCES.keys()
      ])
      )
```

```

        if "rus" in model_key
    ])
)
fi_rus.to_csv(
    os.path.join(output_dir, "rq3-fi-rus.csv"),
    index=False
)
fi_rus

```

```

[12]:

```

	feature	total	1st	2nd	3rd
0	maxNestedBlocks	5	5	0	0
3	loc_method	2	0	1	1
4	uniqueWordsQty_method	2	0	1	1
1	type_anonymous	1	0	1	0
5	wmc_method	1	0	1	0
7	maxNestedBlocksQty	1	0	1	0
2	lcom	1	0	0	1
6	methodsInvokedQty	1	0	0	1
8	cbo_method	1	0	0	1

## 6 RQ 4. How well a model trained with open-source data can generalize to the context of a large-scale enterprise system?

```

[13]: from typing import List

def selected_apache_projects() -> List[str]:
    """
    Returns the name of the selected Apache projects as listed in the "out/
    ↪selection" directory.
    """
    selection_dir = os.path.abspath(os.path.join("out", "selection"))
    return sorted([
        selected.replace(".sh", "")
        for selected in os.listdir(selection_dir)
        if selected.endswith(".sh")
    ])

def load_X_y(project: str):
    dataset_path = os.path.abspath(
        os.path.join("out", "dataset", project, "dataset_full.csv")
    )
    X, y = experiment.load_dataset(
        dataset_path, drops=IGNORED_FEATURES
    )

```

```

    assert X_adyen.shape[1] == X.shape[1]

    return X, y

APACHE_PROJECTS = {
    project: load_X_y(project)
    for project in selected_apache_projects()
}

assert len(APACHE_PROJECTS) == 29

```

```

[14]: for k, v in APACHE_PROJECTS.items():
        print(f"{k:20} {str(v[0].shape):>15}")

```

accumulo	(25458, 63)
ambari	(21997, 63)
archiva	(5995, 63)
bookkeeper	(12711, 63)
cloudstack	(52390, 63)
commons-beanutils	(1176, 63)
cxfs	(33589, 63)
fluo	(2094, 63)
giraph	(8039, 63)
helix	(6790, 63)
ignite	(65181, 63)
jmeter	(8599, 63)
knox	(6821, 63)
lens	(6231, 63)
metamodel	(4122, 63)
myfaces-tobago	(3866, 63)
nutch	(3321, 63)
oodt	(6933, 63)
oozie	(8821, 63)
openmeetings	(4839, 63)
reef	(6150, 63)
sqoop	(3080, 63)
storm	(24208, 63)
syncope	(14915, 63)
tez	(8947, 63)
thrift	(1797, 63)
tomcat	(23793, 63)
zeppelin	(10953, 63)
zookeeper	(5279, 63)



## 6.1 Learning from all Apache projects

```
[15]: X_apache_all = pd.concat(
    [X_apache for X_apache, _ in APACHE_PROJECTS.values()],
    ignore_index=True,
)
y_apache_all = pd.concat(
    [y_apache for _, y_apache in APACHE_PROJECTS.values()],
    ignore_index=True,
)

# Sum of entries must be equals to the number of final entries
assert sum([X.shape[0] for X, _ in APACHE_PROJECTS.values()]) == X_apache_all.
    ↪shape[0]

# apache dataset size, all together
X_apache_all.shape
```

```
[15]: (388095, 63)
```

```
[16]: def rq4():
    scores = []
    model = "rf"
    out = experiment.run(
        model,
        X_train=X_apache_all,
        X_test=X_adyen_test,
        y_train=y_apache_all,
        y_test=y_adyen_test,
        output_to=os.path.join(output_dir, f"rq4-{model}-apache-all.log"),
        tuning_enabled=TUNING_ENABLED
    )
    estimator, score, fi = out
    score["project"] = "apache-all"
    score["training_size"] = X_apache_all.shape[0]
    scores.append(score)

    # Save to the global state this run
    key = f"{model}-apache-all"
    ESTIMATORS[key] = estimator
    FEATURE_IMPORTANCES[key] = fi

    return scores

rq4_scores_all = rq4()
```

## 6.2 Learning from Projects Individually

```
[17]: def rq4_individual():
    scores = []
    model = "rf"
    for project, Xy in APACHE_PROJECTS.items():
        out = experiment.run(
            model,
            X_train=Xy[0].drop(columns=["type"]),
            X_test=X_adyen_test.drop(columns=["type"]),
            y_train=Xy[1].drop(columns=["type"]),
            y_test=y_adyen_test.drop(columns=["type"]),
            output_to=os.path.join(output_dir, f"rq4-{model}-{project}.log"),
            tuning_enabled=TUNING_ENABLED
        )
        estimator, score, fi = out
        score["project"] = project
        score["training_size"] = Xy[0].shape[0]
        scores.append(score)

        # Save to the global state this run
        key = f"{model}-{project}"
        ESTIMATORS[key] = estimator
        FEATURE_IMPORTANCES[key] = fi

    return scores

rq4_scores_individual = rq4_individual()
```

## 6.3 Results

```
[18]: results_rq4 = pd.DataFrame.from_dict(
    merge_scores(
        rq4_scores_all + rq4_scores_individual
    )
)
results_rq4.to_csv(
    os.path.join(output_dir, "rq4.csv"),
    index=False
)
results_rq4.drop(columns=["model", "balancing"]).sort_values(by="prec recall_
↪acc".split(" "), ascending=False)
```

```
[18]:
```

	prec	recall	acc	tn	fp	fn	tp	total	\
0	0.649789	0.259642	0.624013	56497	664	3513	1232	61906	
6	0.643316	0.232244	0.610778	56550	611	3643	1102	61906	

29	0.628972	0.321180	0.652726	56262	899	3221	1524	61906
5	0.621649	0.513172	0.743622	55679	1482	2310	2435	61906
22	0.609648	0.375553	0.677796	56020	1141	2963	1782	61906
19	0.600608	0.208219	0.598363	56504	657	3757	988	61906
1	0.576710	0.263014	0.623494	56245	916	3497	1248	61906
15	0.570383	0.213488	0.600070	56398	763	3732	1013	61906
28	0.566470	0.464278	0.717391	55475	1686	2542	2203	61906
11	0.565574	0.072708	0.534036	56896	265	4400	345	61906
27	0.559633	0.012856	0.506008	57113	48	4684	61	61906
23	0.556629	0.352160	0.664438	55830	1331	3074	1671	61906
17	0.548102	0.404636	0.688471	55578	1583	2825	1920	61906
25	0.539005	0.369863	0.671802	55660	1501	2990	1755	61906
18	0.537485	0.463857	0.715361	55267	1894	2544	2201	61906
24	0.518350	0.288725	0.633227	55888	1273	3375	1370	61906
12	0.515404	0.299684	0.638147	55824	1337	3323	1422	61906
3	0.512434	0.429926	0.697985	55220	1941	2705	2040	61906
13	0.507295	0.190516	0.587578	56283	878	3841	904	61906
7	0.502825	0.093783	0.543043	56721	440	4300	445	61906
2	0.501761	0.240253	0.610225	56029	1132	3605	1140	61906
4	0.482438	0.240253	0.609429	55938	1223	3605	1140	61906
20	0.482411	0.332350	0.651375	55469	1692	3168	1577	61906
21	0.481291	0.295469	0.634517	55650	1511	3343	1402	61906
8	0.466184	0.122023	0.555212	56498	663	4166	579	61906
10	0.464710	0.403793	0.682592	54954	2207	2829	1916	61906
14	0.463646	0.185458	0.583825	56143	1018	3865	880	61906
9	0.433325	0.356797	0.659032	54947	2214	3052	1693	61906
16	0.381102	0.255005	0.610314	55196	1965	3535	1210	61906
26	0.242255	0.201054	0.574425	54177	2984	3791	954	61906

	mean_fit_time	std_fit_time	mean_test_score	std_test_score	\
0	58.143208	2.070854	0.631243	0.019912	
6	0.203064	0.005256	0.645367	0.092532	
29	2.783138	0.063516	0.658597	0.047932	
5	33.432431	0.457052	0.701172	0.049714	
22	0.309914	0.005822	0.637366	0.061589	
19	0.497548	0.021913	0.587284	0.042232	
1	1.674008	0.048119	0.659687	0.063352	
15	0.136482	0.003493	0.580599	0.052718	
28	5.079360	0.283219	0.608371	0.052231	
11	4.917649	0.173595	0.604817	0.045741	
27	12.193701	0.820597	0.578226	0.042667	
23	1.273365	0.089453	0.614461	0.034391	
17	1.488858	0.051147	0.743647	0.036559	
25	4.055725	0.077063	0.689929	0.028932	
18	2.839530	0.041913	0.687929	0.051630	
24	6.344290	0.078909	0.689769	0.046570	
12	3.994664	0.038367	0.682666	0.026414	

3	2.384433	0.067282	0.575410	0.027376
13	2.780408	0.077759	0.581728	0.040740
7	2.043416	0.086845	0.555690	0.021273
2	10.048509	0.087201	0.618623	0.044476
4	5.754161	0.115014	0.652809	0.040228
20	0.247272	0.020979	0.581521	0.046188
21	2.261823	0.044616	0.745460	0.040013
8	0.873031	0.058591	0.498227	0.113801
10	3.026489	0.033126	0.610663	0.069676
14	0.319387	0.027164	0.596098	0.041634
9	0.345628	0.017250	0.696688	0.048262
16	0.139436	0.003109	0.652011	0.051416
26	0.634905	0.069358	0.587256	0.100076

	project	training_size
0	apache-all	388095
6	commons-beanutils	1176
29	zookeeper	5279
5	cloudstack	52390
22	sqoop	3080
19	oozie	8821
1	accumulo	25458
15	metamodel	4122
28	zeppelin	10953
11	ignite	65181
27	tomcat	23793
23	storm	24208
17	nutch	3321
25	tez	8947
18	oodt	6933
24	syncope	14915
12	jmeter	8599
3	archiva	5995
13	knox	6821
7	cxfr	33589
2	ambari	21997
4	bookkeeper	12711
20	openmeetings	4839
21	reef	6150
8	fluo	2094
10	helix	6790
14	lens	6231
9	giraph	8039
16	myfaces-tobago	3866
26	thrift	1797